

Parallel Algorithms for Graph Similarity and Matching

*Advanced Algorithms and Parallel
Programming Course Project
Marco Bacis, Antonio Di Bello*

Outline


- Problem Overview
- Network Similarity Decomposition
- Auction-based matching
- Experimental Set and Results

Outline

- Problem Overview
- Network Similarity Decomposition
- Auction-based matching
- Experimental Set and Results

Problem Definition

- Given two graphs:
 - How similar is each vertex in the first graph to each vertex in the second?
 - What is the best match for each vertex in the first graph to a vertex in the second graph?




Graph Similarity

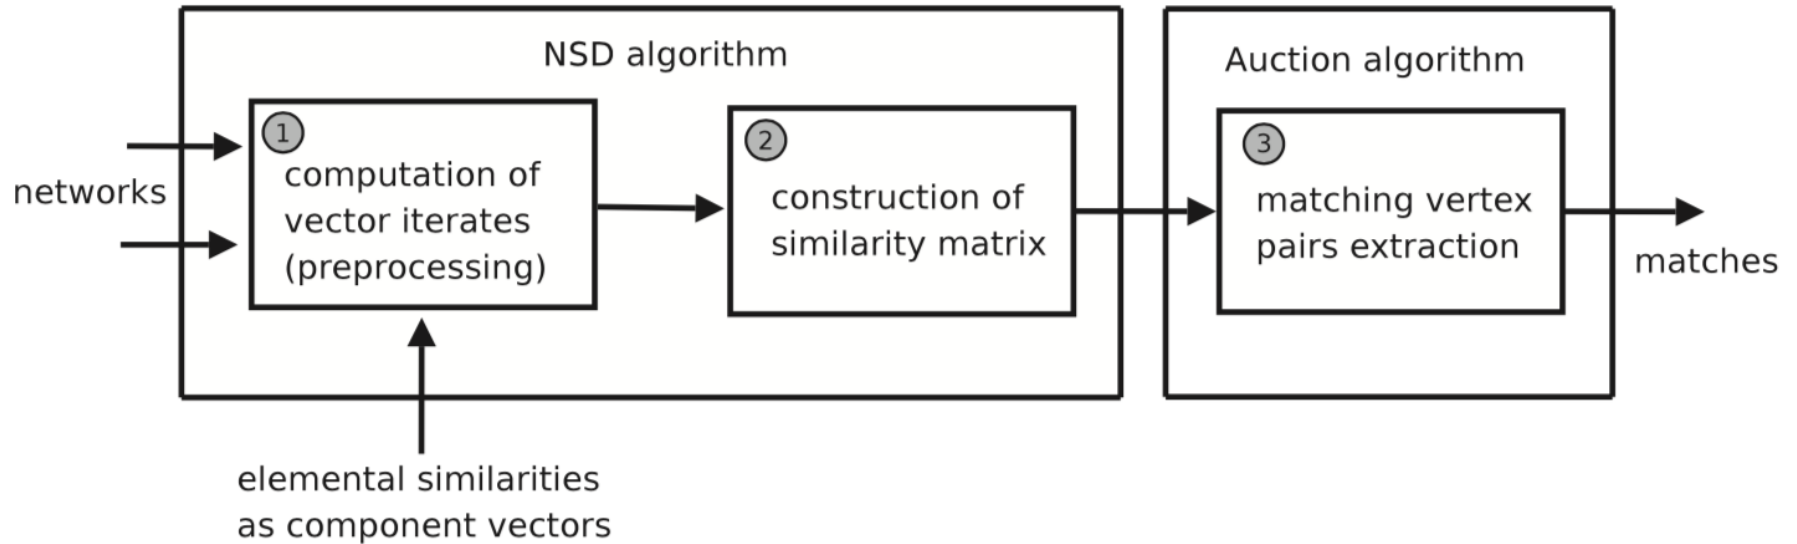
- Two main categories:
 - Single/ Global similarity score (scalar)
 - Vertex-wise similarity score (matrix)
- Approaches:
 - GRAAL family, the “seed and extend” idea
 - IsoRank, vertex similarity scores using vertex attributes and topological similarities
 - NSD, low-rank decompositions of the matrix to decouple its construction process

Bipartite Graph Matching

- Matching M of vertices over weighted edge :
 - a vertex is an endpoint of at most one matching edge
 - sum over the matched edges is maximized
- Implementations:
 - Augmenting path
 - Hungarian method
 - Auction-based algorithm



Overview



- Network Similarity Decomposition (NSD) matrix computation
- Auction-Based matching using the similarity matrix
- Integrated approach (NSD + Auction on the same processes)

Outline

- Problem Overview
- Network Similarity Decomposition
- Auction-based matching
- Experimental Set and Results

Network Similarity Decomposition (NSD)

- Iterative definition of X :

$$X \leftarrow \alpha \tilde{B} X \tilde{A}^T + (1 - \alpha) H$$

where:

- X = Resulting similarity matrix
- \tilde{A}, \tilde{B} = Normalized/Transposed Adjacency matrices
- α = scale factor
- H = a-priori vertex similarity matrix

Network Similarity Decomposition (NSD)

NSD relies on low-rank representations of the H matrix, into a sum of outer products of vectors.

- **Singular Value Decomposition (SVD)** decompose H into eigenvalue and pair of eigenvector .
- Other possible methods:
 - Non-negative Matrix Factorization (NMF)
 - or other decomposition method

This enables further possibilities of parallelization.

$$H = \sum_{i=1}^r \sigma_i u_i v_i^T$$


NSD Algorithm Pseudocode

Algorithm 1 NSD: Calculate $X^{(n)}$ given $A, B, \{w_i, z_i | i = 1, \dots, s\}, \alpha$ and n

```
1: compute  $\tilde{A}, \tilde{B}$ 
2: for  $i = 1$  to  $s$  do
3:    $w_i^{(0)} \leftarrow w_i$ 
4:    $z_i^{(0)} \leftarrow z_i$ 
5:   for  $k = 1$  to  $n$  do
6:      $w_i^{(k)} \leftarrow \tilde{B}w_i^{(k-1)}$ 
7:      $z_i^{(k)} \leftarrow \tilde{A}z_i^{(k-1)}$ 
8:   end for
9:   zero  $X_i^{(n)}$ 
10:  for  $k = 0$  to  $n - 1$  do
11:     $X_i^{(n)} \leftarrow X_i^{(n)} + \alpha^k w_i^{(k)} z_i^{(k)T}$ 
12:  end for
13:   $X_i^{(n)} \leftarrow (1 - \alpha)X_i^{(n)} + \alpha^n w_i^{(n)} z_i^{(n)T}$ 
14: end for
15:  $X^{(n)} \leftarrow \sum_{i=1}^s X_i^{(n)}$ 
```

NSD Parallel Algorithm

- Centralized (parallel) computation of w, z
- Distributed X computation
 - w split over a group of processes
 - Partial computation on each node/process



Outline

- Problem Overview
- Network Similarity Decomposition
- Auction-based matching
- Experimental Results

Auction-based graph matching

- V_a and V_b represent the set of buyers and objects where $n_a \leq n_b$

The algorithm consists of three phases:

- the **initialization** phase
- the **bidding** phase
- the **assignment** phase

Auction-based graph matching

Algorithm 2 Sequential Auction Algorithm for Maximum Weighted Matching

Input: Bipartite graph $G = (V_A, V_B, E, w)$

Output: Matching M



```
1:  $M \leftarrow \emptyset$  ▷ current matching
2:  $I \leftarrow \{i : 1 \leq i \leq n_A\}$  ▷ set of unassigned buyers
3:  $p_j \leftarrow 0$  for  $j = 1, \dots, n_B$  ▷ initialize prices for objects
4: initialize( $\varepsilon$ ) ▷ initialize  $\varepsilon$ 
5: while  $I \neq \emptyset$  do ▷ auction iteration
6:    $j_i \leftarrow \arg \max_j \{x_{ij} - p_j\}$  ▷ find best object of buyer  $i$ 
7:    $u_i \leftarrow x_{ij_i} - p_{j_i}$  ▷ store profit of the most valuable object
8:    $v_i \leftarrow \max_{j \neq j_i} \{x_{ij} - p_j\}$  ▷ store second-best profit
9:    $p_{j_i} \leftarrow p_{j_i} + u_i - v_i + \varepsilon$  ▷ update price with the bid  $u_i - v_i$  and  $\varepsilon$ 
10:   $M \leftarrow M \cup \{i, j_i\}; I \leftarrow I \setminus \{i\}$  ▷ assign buyer to the desired object
11:   $M \leftarrow M \setminus \{k, j_i\}; I \leftarrow I \cup \{k\}$  ▷ free previous owner  $k$  if available
12:  update( $\varepsilon$ ) ▷ increment/decrement  $\varepsilon$ 
13: end while
```

Auction-based graph matching

Detected parallelism:

- Bids of free buyers simultaneously computed:
 - Each free buyer computes a bid for the most-valuable object according to the current price
 - The prices of the objects are updated according to the highest bids
 - Exchanging through messages only locally altered prices
- 1D row-wise distribution of the similarity matrix to facilitate buyers partitioning

Auction-based graph matching



- Local auction on each process (local free buyers)
- Global check and price/free buyers update
- Convergence when no free buyers left

Outline

- Problem Overview
- Network Similarity Decomposition
- Auction-based matching
- Experimental Set and Results

Implementation

- Hybrid parallel programming model
 - OpenMP for MVP, similarity matrix and local auction
 - MPI for problem partitioning (each task works on a subset of the matrix/auction)
- Tested on a single dual-socket machine
 - Ideally, one mpi task for each socket
 - Maximize number of openmp threads on each socket

Complete Algorithm

Algorithm 6 NSD-based Parallel Graph Matching

- 1: $\square = \text{root process, no labels} = \text{all processes } r$
 - 2: \square load adjacency matrices A, B and component vectors w_i, z_i ;
 - 3: \square compute \tilde{A}, \tilde{B} ;
 - 4: **broadcast** \tilde{A}, w_i, z_i ;
 - 5: distribute \tilde{B} by row blocks \triangleright each process r gets its \tilde{B}_r part;
 - 6: for all components i and steps k ($z_i^{(0)} = z_i, w_i^{(0)} = w_i$)
 - 7: compute vector iterates $z_i^{(k)} \leftarrow \tilde{A}z_i^{(k-1)}$
 - 8: compute vector iterates $w_{i,r}^{(k)} \leftarrow \tilde{B}_r w_i^{(k-1)}$, **gather** $w_i^{(k)}$ (*// matvec*);
 - 9: compute *row-wise* the local similarity matrix X_r (*embarrassingly //*)
 - 10: \triangleright NSD-based, *sparsify* if needed (sort row entries, keep largest ones);
 - 11: compute weighted matchings by *// auction*
 - 12: \triangleright matching permutation lands on root;
 - 13: \square compute number of conserved edges, similarity rate;
-

Complete Algorithm

Algorithm 6 NSD-based Parallel Graph Matching

- 1: $\square = \text{root process, no labels} = \text{all processes } r$
- 2: \square load adjacency matrices A, B and component vectors w_i, z_i ;
- 3: \square compute \tilde{A}, \tilde{B} ;

- Sparse representation of \tilde{A}, \tilde{B}
- Fast transpose/normalization thanks to low number of nonzero values

Complete Algorithm

- 4: broadcast \tilde{A} , w_i , z_i ;
- 5: distribute \tilde{B} by row blocks ▷ each process r gets its \tilde{B}_r part;
- 6: for all components i and steps k ($z_i^{(0)} = z_i$, $w_i^{(0)} = w_i$)
- 7: compute vector iterates $z_i^{(k)} \leftarrow \tilde{A}z_i^{(k-1)}$
- 8: compute vector iterates $w_{i,r}^{(k)} \leftarrow \tilde{B}_r w_i^{(k-1)}$, **gather** $w_i^{(k)}$ (*// matvec*);
- 9: compute *row-wise* the local similarity matrix X_r (*embarrassingly //*)

- Sparse matrix broadcast/allgather
- OpenMP for Matrix-Vector product and for X iterations

X Iterations

```
192     for (int i = 1; i < n; i++) {
193         vector_t W_gathered;
194
195         //std::cout << "Worker " << rank << " iterate " << i << std::endl;
196         if (i == 1) W_gathered = W;
197         else W_gathered = allgather_vector(W_i[i-1]);
198
199         W_i[i] = matvect_prod(B, W_gathered);
200         Z_i[i] = matvect_prod(A, Z_i[i-1]);
201     }
202
203     #pragma omp parallel for shared(X, W_i, Z_i)
204     for(unsigned int y = 0; y < X.size1(); y++) {
205         float alpha_pow = alpha;
206         //if (y%1000 == 0) std::cout << "Worker " << rank << " row " << y << std::endl;
207         //#pragma omp parallel for collapse(2) private(i,x,alpha_pow)
208         for(int i = 0; i < n-1; i++) {
209             for(unsigned int x = 0; x < X.size2(); x++) {
210                 X(y,x) += alpha_pow * W_i[i][y] * Z_i[i][x];
211             }
212             alpha_pow *= alpha;
213         }
214         //last step (n)
215         //#pragma omp parallel for
216         for(unsigned int x = 0; x < X.size2(); x++)
217             X(y,x) = (1-alpha) * X(y,x) + alpha_pow * W_i[n-1][y] * Z_i[n-1][x];
218     }
```

Complete Algorithm

11: compute weighted matchings by `// auction`

12: ▷ matching permutation lands on root;

- Parallel argmax (OpenMP custom reduction)
- Randomized epsilon scaling
- Allgather for local auctions
results (price and assignement) merge

Parallel argmax

```
123     struct BidResult res = {.obj=-1, .buyer=-1, .maxP=-FLT_MAX, .secondP=-FLT_MAX};
124
125     #pragma omp parallel
126     {
127     #pragma omp for collapse(2) reduction(bidReduce:res)
128     for (unsigned int i = 0; i < freeBuyer.size(); i++) {
129         for (int j=0; j<nb; j++) {
130             auto it = freeBuyer.begin();
131             advance(it,i);
132             if (X(*it,j) - price[j] > res.maxP) {
133                 if (j != res.obj) {
134                     res.secondP = res.maxP;
135                 }
136                 res.obj = j;
137                 res.maxP = X(*it,j) - price[j];
138                 res.buyer = *it;
139             } else if (j!=res.obj && X(*it,j) - price[j] > res.secondP) {
140                 res.secondP = X(*it,j) - price[j];
141             }
142         }
143     }
144 }
```

Results exchange


```
157      /* Gather changed prices */
158      sendBuyer.clear();
159      sendPrice[0] = localPrice; //price
160      sendPrice[1] = res.obj; // object
161      MPI_Allgather(sendPrice,2,MPI_FLOAT,receivedPrices,2,MPI_FLOAT,MPI_COMM_WORLD);
162      for (int i=0; i<worldSize*2; i+=2) {
163          if (price[receivedPrices[i+1]] < receivedPrices[i]) {
164              price[receivedPrices[i+1]] = receivedPrices[i];
165              //se qualcuno dei miei local aveva comprato quell'oggetto ora lo perde.
166              if (assigned[receivedPrices[i+1]] != -1) {
167
168                  if (debug)
169                      cout << worldRank << " : localbuyer " << assigned[receivedPrices[i+1]] << " lose obj " << receivedPrices[i+1]
170
171                      sendBuyer.push_back(assigned[receivedPrices[i+1]] + (worldRank*(na/worldSize)) + 1);
172
173                      match[assigned[receivedPrices[i+1]]] = -1;
174                      freeBuyer.insert(assigned[receivedPrices[i+1]]);
175                      assigned[receivedPrices[i+1]] = -1;
176
177              }
178          } else if (price[receivedPrices[i+1]] == receivedPrices[i] && receivedPrices[i+1] == res.obj) {
179              //se qualcuno ha offerto lo stesso prezzo del mio vincitore local e il suo indice globale è minore del mio, perdo l'ogg
180              localPrice--;
181              }
182      }
183  }
```

Evaluation

- Graphs tested
 - Yeast protein: 2361 nodes
 - Wikivote graph: 8297 nodes
- Matching with subgraphs (and self)
 - Yeast -> 500/1000 nodes subgraph
 - Wikivote -> 1000 nodes subgraph

Results


Yeast (complete) vs Yeast subgraph (500 nodes)



Only 1 MPI Task

Results

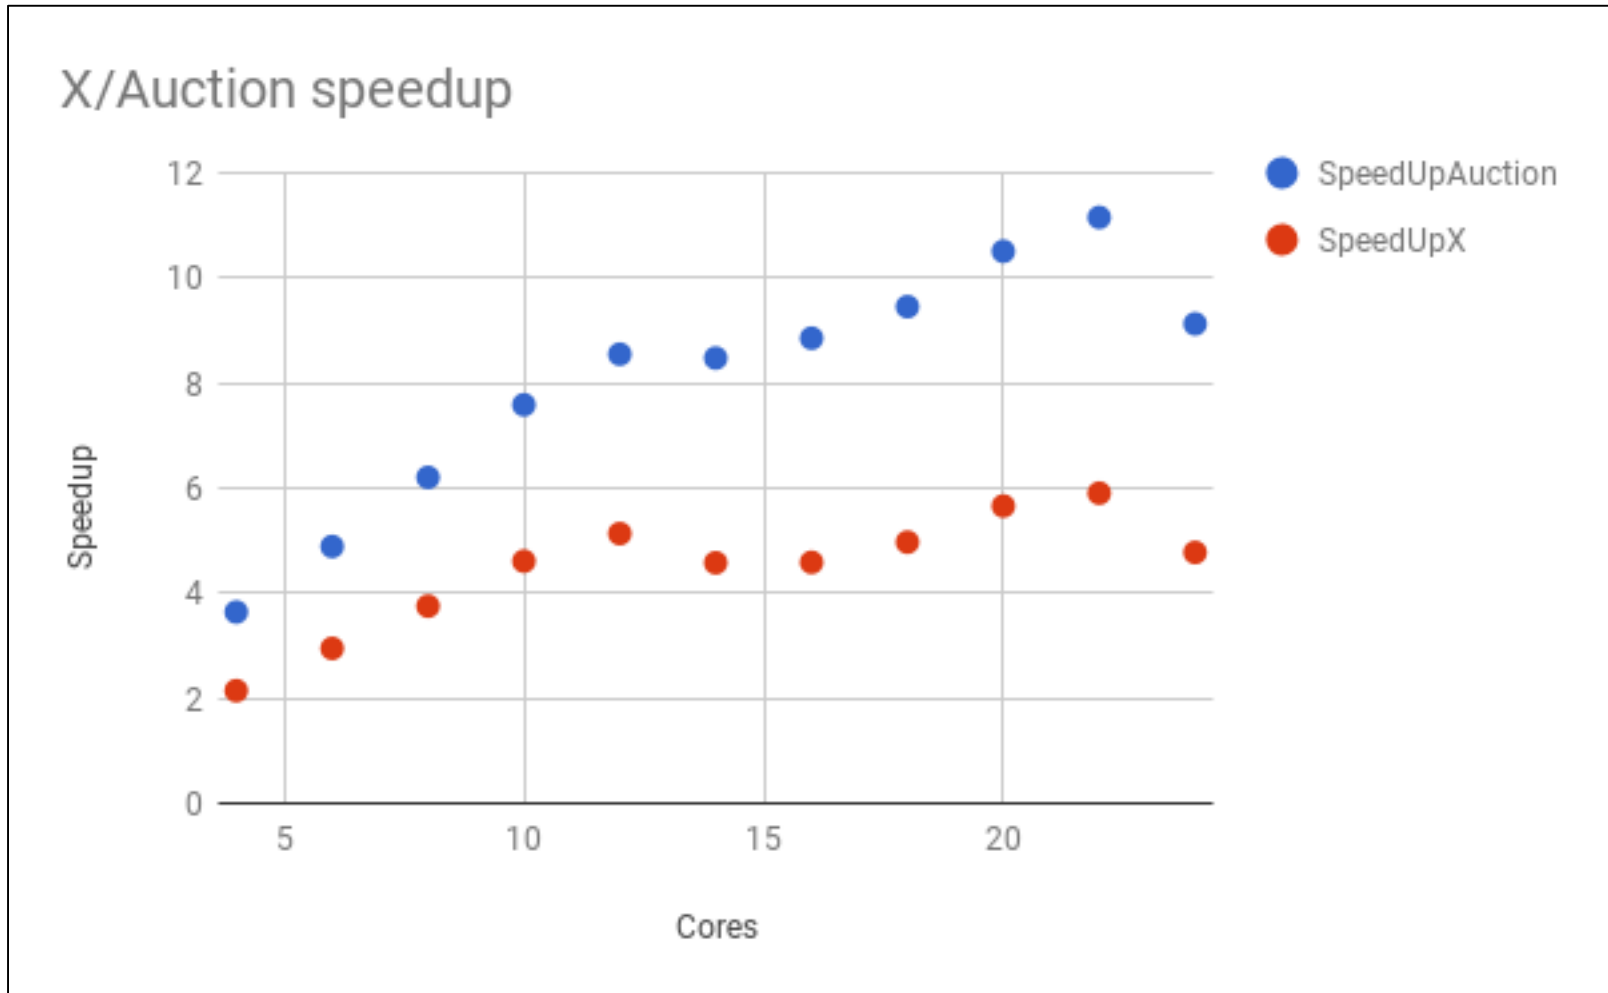
Yeast (complete) vs Yeast subgraph (500 nodes)



2 MPI Tasks

Results


Yeast (complete) vs Yeast subgraph (500 nodes)



2 MPI Tasks

Results


Wikivote (complete) vs Wikivote subgraph (1000 nodes)



1 MPI Task

Results

Wikivote (complete) vs Wikivote subgraph (1000 nodes)



2 MPI Tasks

Conclusions

- 2 dimensions scaling (MPI tasks – OMP threads)
- Linear speedup w.r.t number of cores used
- Slow auction convergence due to price wars
- Using 2 MPI tasks leads to best results
(as expected)

Question

